

Multiple Linear Regression

Machine Learning

Dr. Adnan Abid

Regressions

Simple Linear Regression

$$y = b_0 + b_1 * x_1$$

Multiple Linear Regression

Dependent variable (DV)

Independent variables (IVs)


$$y = b_0 + b_1 * x_1 + b_2 * x_2 + \dots + b_n * x_n$$

Dummy Variables

Profit	R&D Spend	Admin	Marketing	State
192,261.83	165,349.20	136,897.80	471,784.10	New York
191,792.06	162,597.70	151,377.59	443,898.53	California
191,050.39	153,441.51	101,145.55	407,934.54	California
182,901.99	144,372.41	118,671.85	383,199.62	New York
166,187.94	142,107.34	91,391.77	366,168.42	California

$$y = b_0 + b_1 * x_1 + b_2 * x_2 + b_3 * x_3 + ???$$

Dummy Variables

Profit	R&D Spend	Admin	Marketing	State
192,261.83	165,349.20	136,897.80	471,784.10	New York
191,792.06	162,597.70	151,377.59	443,898.53	California
191,050.39	153,441.51	101,145.55	407,934.54	California
182,901.99	144,372.41	118,671.85	383,199.62	New York
166,187.94	142,107.34	91,391.77	366,168.42	California

New York	California

$$y = b_0 + b_1 * x_1 + b_2 * x_2 + b_3 * x_3 + ???$$

Dummy Variables

Profit	R&D Spend	Admin	Marketing	State	New York	California
192,261.83	165,349.20	136,897.80	471,784.10	New York	1	0
191,792.06	162,597.70	151,377.59	443,898.53	California	0	1
191,050.39	153,441.51	101,145.55	407,934.54	California	0	1
182,901.99	144,372.41	118,671.85	383,199.62	New York	1	0
166,187.94	142,107.34	91,391.77	366,168.42	California	0	1

$$y = b_0 + b_1 * x_1 + b_2 * x_2 + b_3 * x_3 + ???$$

Dummy Variables

Profit	R&D Spend	Admin	Marketing	State
192,261.83	165,349.20	136,897.80	471,784.10	New York
191,792.06	162,597.70	151,377.59	443,898.53	California
191,050.39	153,441.51	101,145.55	407,934.54	California
182,901.99	144,372.41	118,671.85	383,199.62	New York
166,187.94	142,107.34	91,391.77	366,168.42	California

Dummy Variables

New York	California
1	0
0	1
0	1
1	0
0	1

$$y = b_0 + b_1 * x_1 + b_2 * x_2 + b_3 * x_3 + b_4 * D_1$$

1. Add dummy variables for all possible values in the attribute with categorical data, which is STATE in this example.
2. Populate them with One Hot Encoding method
3. Replace the column with Dummy variables

Dummy Variables

Profit	R&D Spend	Admin	Marketing	State
192,261.83	165,349.20	136,897.80	471,784.10	New York
191,792.06	162,597.70	151,377.59	443,898.53	California
191,050.39	153,441.51	101,145.55	407,934.54	California
182,901.99	144,372.41	118,671.85	383,199.62	New York
166,187.94	142,107.34	91,391.77	366,168.42	California

Dummy Variables

New York	California
1	0
0	1
0	1
1	0
0	1

$$y = b_0 + b_1 * x_1 + b_2 * x_2 + b_3 * x_3 + b_4 * D_1$$

1. Add dummy variables for all possible values in the attribute with categorical data, which is STATE in this example.
2. Populate them with One Hot Encoding method
3. Replace the column with Dummy variables

1. D1 works like a switch, when it is 1 then it means New York and in case of 0 it is California
2. Thus we may not include the last one, as the information will be complete even when we eliminate it.-

Dummy Variables

Profit	R&D Spend	Admin	Marketing	State
192,261.83	165,349.20	136,897.80	471,784.10	New York
191,792.06	162,597.70	151,377.59	443,898.53	California
191,050.39	153,441.51	101,145.55	407,934.54	California
182,901.99	144,372.41	118,671.85	383,199.62	New York
166,187.94	142,107.34	91,391.77	366,168.42	California

Dummy Variables

New York	California
1	0
0	1
0	1
1	0
0	1

$$y = b_0 + b_1 * x_1 + b_2 * x_2 + b_3 * x_3 + b_4 * D_1$$



1. D1 works like a switch, when it is 1 then it means New York and in case of 0 it is California
2. Thus we may not include the last one, as the information will be complete even when we eliminate it.

1. Theoretically, we SHOULD NOT include all the dummy variables in the equation.
2. Intuition is that when D1 is 0, then $b_4 * D_1$ is also 0, and we may think that there is no coefficient for California.
3. Well, the coefficient for California is going to be included in b_0 , i.e. when D1 is 0 the equation becomes an expression for California.
4. When D1 is 1 then the expression becomes the difference between New York and California

Dummy Variable Trap

Profit	R&D Spend	Admin	Marketing	State
192,261.83	165,349.20	136,897.80	471,784.10	New York
191,792.06	162,597.70	151,377.59	443,898.53	California
191,050.39	153,441.51	101,145.55	407,934.54	California
182,901.99	144,372.41	118,671.85	383,199.62	New York
166,187.94	142,107.34	91,391.77	366,168.42	California

Dummy Variables

New York	California
1	0
0	1
0	1
1	0
0	1

$$y = b_0 + b_1 * x_1 + b_2 * x_2 + b_3 * x_3$$



$$+ b_4 * D_1 + \underline{b_5 * D_2}$$



Activate Windows
Go to Settings to activate Windows.

1. The fact that one or more independent variables in a linear regression predict another is called multi-collinearity.
2. If so, the model cannot distinguish the facts of D1 from the facts of D2.
3. This is called Dummy Variable Trap.
4. Thus, you cannot have b_0 , D_1 , and D_2 in the model at the same time.

Dummy Variable Trap

Profit	R&D Spend	Admin	Marketing	State	Dummy Variables	
					New York	California
192,261.83	165,349.20	136,897.80	471,784.10	New York	1	0
191,792.06	162,597.70	151,377.59	443,898.53	California	0	1
191,050.39	153,441.51	101,145.55	407,934.54	California	0	1
182,901.99	144,372.41	118,671.85	383,199.62	New York	1	0
166,187.94	142,107.34	91,391.77	366,168.42	California	0	1

$$y = b_0 + b_1 * x_1 + b_2 * x_2 + b_3 * x_3 + b_4 * D_1 + \underline{b_5 * D_2}$$

1. Never include all the dummy variables.
2. Always ignore 1 dummy variable. If you have 10 dummy variables, then include 9 and leave 1; if you have 100 dummy variables include 99 and leave 1.
3. Never include all the dummy variables.
2. If you have more than one categorical data attributes, then you have to do the same for the other variable too, i.e. include n-1 dummy variables and leave one out.
3. For instance, if you have another variable that tell about the target sector of the organization, then we shall follow the same process all over for this variable too.

Dummy Variable Trap

Profit	R&D Spend	Admin	Marketing	State
192,261.83	165,349.20	136,897.80	471,784.10	New York
191,792.06	162,597.70	151,377.59	443,898.53	California
191,050.39	153,441.51	101,145.55	407,934.54	California
182,901.99	144,372.41	118,671.85	383,199.62	New York
166,187.94	142,107.34	91,391.77	366,168.42	California

Dummy Variables

New York	California
1	0
0	1
0	1
1	0
0	1

$$y = b_0 + b_1 * x_1 + b_2 * x_2 + b_3 * x_3$$

$$+ b_4 * D_1 + \cancel{b_5 * D_2}$$


Always omit one
dummy variable

Building A Model

(Step-By-Step)

Activate Windows
Go to Settings to activate Windows.

Building A Model



1. There are a lot of variables
2. Should we include all the independent variables?
3. In fact, we need to include the ones which are helpful in prediction, and drop the rest.
4. There are many methods for doing it...

Building A Model

5 methods of building models:

1. All-in
2. Backward Elimination
3. Forward Selection
4. Bidirectional Elimination
5. Score Comparison


Building A Model

“All-in” – cases:

- Prior knowledge; OR
- You have to; OR
- Preparing for Backward Elimination



Building A Model



Backward Elimination

STEP 1: Select a significance level to stay in the model (e.g. $SL = 0.05$)



STEP 2: Fit the full model with all possible predictors



STEP 3: Consider the predictor with the highest P-value. If $P > SL$, go to STEP 4, otherwise go to FIN



STEP 4: Remove the predictor




STEP 5: Fit model without this variable*



1. Remove the attribute with highest P-value, if the P-value is greater than the threshold Significance Level.
2. Then go to Step 3 and rebuild the model
3. If there is not such variable with P-value higher than SL, the go to FINISH (FIN) state, which means the model is ready and all the remain attributes are statistically significant for prediction.

Building A Model

Forward Selection



STEP 1: Select a significance level to enter the model (e.g. $SL = 0.05$)



STEP 2: Fit all simple regression models $y \sim x_n$. Select the one with the lowest P-value



STEP 3: Keep this variable and fit all possible models with one extra predictor added to the one(s) you already have



STEP 4: Consider the predictor with the lowest P-value. If $P < SL$, go to STEP 3, otherwise go to FIN

1. Add one variable at a time.
2. Looks reverse of backward elimination, but is much more tedious and complex.
3. Start with a single variable (Simple Regression) like manner.
4. Choose the variable which has lowest SL.
5. Then add each of the remaining variables one by one; and choose the one with lowest SL value.
6. Keep doing it till you find variable with P values $< SL$.
7. If you are unable to find any variable which, if added, to the model does not offer $PL < SL$.

Building A Model

Forward Selection

STEP 1: Select a significance level to enter the model (e.g. $SL = 0.05$)



STEP 2: Fit all simple regression models $y \sim x_n$. Select the one with the lowest P-value



STEP 3: Keep this variable and fit all possible models with one extra predictor added to the one(s) you already have




STEP 4: Consider the predictor with the lowest P-value. If $P < SL$, go to STEP 3, otherwise go to FIN

1. Stop when you are unable to find a variable where $P < SL$
2. IMPORTANT: Use the previous model, as this attribute does not satisfy the SL criteria. Thus this model should not be included in the model.




FIN: Keep the previous model



Building A Model

Bidirectional Elimination

STEP 1: Select a significance level to enter and to stay in the model
e.g.: SLENTER = 0.05, SLSTAY = 0.05



STEP 2: Perform the next step of Forward Selection (new variables must have: $P < \text{SLENTER}$ to enter)

STEP 3: Perform ALL steps of Backward Elimination (old variables must have $P < \text{SLSTAY}$ to stay)

STEP 4: No new variables can enter and no old variables can exit

FIN: Your Model Is Ready

Activate Windows
Go to Settings to activate Windows.

Building A Model

All Possible Models

STEP 1: Select a criterion of goodness of fit (e.g. Akaike criterion)



STEP 2: Construct All Possible Regression Models: $2^N - 1$ total combinations



STEP 3: Select the one with the best criterion



Building A Model

All Possible Models

STEP 1: Select a criterion of goodness of fit (e.g. Akaike criterion)



STEP 2: Construct All Possible Regression Models: $2^N - 1$ total combinations



STEP 3: Select the one with the best criterion



FIN: Your Model Is Ready



Example:
10 columns means
1,023 models

Activate Windows
Go to Settings to activate Windows.

Building A Model

5 methods of building models:

1. All-in
2. Backward Elimination
3. Forward Selection
4. Bidirectional Elimination
5. Score Comparison

Console 1/A

Date: Sun, 25 Sep 2016 Time: 16:13:59

Prob (F-statistic): 1.34e-27

Log-Likelihood: -525.38

No. Observations: 50

AIC: 1063.

Df Residuals: 44

BIC: 1074.

Df Model: 5

Covariance Type: nonrobust

	coef	std err	t	P> t	[95.0% Conf. Int.]	
const	5.013e+04	6884.820	7.281	0.000	3.62e+04	6.4e+04
x1	198.7888	371.007	0.059	0.953	-6595.030	6992.607
x2	-41.8870	3256.039	-0.013	0.990	-6604.003	6520.229
x3	0.8060	0.046	17.369	0.000	0.712	0.900
x4	-0.0270	0.052	-0.517	0.608	-0.132	0.078
x5	0.0270	0.017	1.574	0.123	-0.008	0.062

Omnibus: 14.782 Durbin-Watson: 1.283

Prob(Omnibus): 0.001 Jarque-Bera (JB): 21.266

Permissions: RW End-of-lines: LF Encoding: UTF-8-GUESSED Line: 49 Column: 24 Memory: 55 %

IPython console

Console 1/A

	coef	std err	t	P> t	[95.0% Conf. Int.]	
const	5.011e+04	6647.870	7.537	0.000	3.67e+04	6.35e+04
x1	220.1585	2900.536	0.076	0.940	-5621.821	6062.138
x2	0.8060	0.046	17.606	0.000	0.714	0.898
x3	-0.0270	0.052	-0.523	0.604	-0.131	0.077
x4	0.0270	0.017	1.592	0.118	-0.007	0.061

Omnibus: 14.758 Durbin-Watson: 1.282

Prob(Omnibus): 0.001 Jarque-Bera (JB): 21.172

Skew: -0.948 Prob(JB): 2.53e-05

Kurtosis: 5.563 Cond. No. 1.40e+06

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Permissions: RW End-of-lines: LF Encoding: UTF-8-GUESSED Line: 50 Column: 17 Memory: 56 %

```

41# Building the optimal model using Backward Elimination
42import statsmodels.formula.api as sm
43X = np.append(arr = np.ones((50, 1)).astype(int), values = X, axis =
44X_opt = X[:, [0, 1, 2, 3, 4, 5]]
45regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()
46regressor_OLS.summary()
47X_opt = X[:, [0, 1, 3, 4, 5]]
48regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()
49regressor_OLS.summary()
50X_opt = X[:, [0, 3, 4, 5]]
51regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()
52regressor_OLS.summary()

```


Spyder (Python 3.5)

/Users/Hadelin/Desktop/Machine Learning A-Z/Part 2 - Regression/Section 5 - Multiple Linear Regression

Editor - /Users/Hadelin/Desktop/Machine Learning A-Z/Part 2 - Regression/Section 5 - Multiple Linear Regression/multiple_linear_regression.py

data_preprocessing_template.py

multiple_linear_regression.py*

```
18 X = onehotencoder.fit_transform(X).toarray()
19
20 # Avoiding the Dummy Variable Trap
21 X = X[:, 1:]
22
23 # Splitting the dataset into the Training set and Test set
24 from sklearn.cross_validation import train_test_split
25 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_s
26
27 # Feature Scaling
28 """from sklearn.preprocessing import StandardScaler
29 sc_X = StandardScaler()
30 X_train = sc_X.fit_transform(X_train)
31 X_test = sc_X.transform(X_test)"""
32
33 # Fitting Multiple Linear Regression to the Training set
34 from sklearn.linear_model import LinearRegression
35 regressor = LinearRegression()
36 regressor.fit(X_train, y_train)
37
38 # Predicting the Test set results
39 y_pred = regressor.predict(X_test)
40
41 # Building the optimal model using Backward Elimination
42 import statsmodels.formula.api as sm
43 X = np.append(arr = np.ones((50, 1)).astype(int), values = X, axis = 1)
44 X_opt = X[:, [0, 1, 2, 3, 4, 5]]
45 regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()
46 regressor_OLS.summary()
47 X_opt = X[:, [0, 1, 3, 4, 5]]
48 regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()
49 regressor_OLS.summary()
50 X_opt = X[:, [0, 3, 4, 5]]
51 regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()
52 regressor_OLS.summary()
53 X_opt = X[:, [0, 3, 5]]
54 regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()
55 regressor_OLS.summary()
56 X_opt = X[:, [0, 3]]
57 regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()
58 regressor_OLS.summary()
```

Variable explorer

Name	Type	Size	Value
X	float64	(50, 6)	array([[1.00000000e+00, 0.00000000e+00, 1.00000000e+00, 1.65349200e+05, 1.36897800e+05, 4.71784100e+05],
X_opt	float64	(50, 2)	array([[1.00000000e+00, 1.65349200e+05],
X_test	float64	(10, 5)	array([[1.00000000e+00, 0.00000000e+00, 6.60515200e+04,
X_train	float64	(40, 5)	array([[1.00000000e+00, 0.00000000e+00, 5.54939500e+04,
dataset	DataFrame	(50, 5)	Column names: R&D Spend, Administration, Marketing Spend, State, Profit
y	float64	(50,)	array([192261.83, 191792.06, 191050.39, 182901.99, 166187.94,
y_pred	float64	(10,)	array([103015.20159796, 132582.27760815, 132447.73845175,
y_test	float64	(10,)	array([103282.38, 144259.4 , 146121.95, 77798.83, 191050.39,
y_train	float64	(40,)	array([96778.92, 96479.51, 105733.54, 96712.8 , 124266.9 ,

Object inspector

Variable explorer

File explorer

iPython console

Console 1/A

```
Dep. Variable: y R-squared: 0.947
Model: OLS Adj. R-squared: 0.945
Method: Least Squares F-statistic: 849.8
Date: Sun, 25 Sep 2016 Prob (F-statistic): 3.50e-32
Time: 19:17:55 Log-Likelihood: -527.44
No. Observations: 50 AIC: 1059.
Df Residuals: 48 BIC: 1063.
Df Model: 1
Covariance Type: nonrobust

=====
coef std err t P>|t| [95.0% Conf. Int.]
-----
const 4.903e+04 2537.897 19.320 0.000 4.39e+04 5.41e+04
x1 0.8543 0.029 29.151 0.000 0.795 0.913
=====
Omnibus: 13.727 Durbin-Watson: 1.116
Prob(Omnibus): 0.001 Jarque-Bera (JB): 18.536
```


Polynomial Linear Regression

Regressions

Simple
Linear
Regression

$$y = b_0 + b_1 x_1$$


Multiple
Linear
Regression

$$y = b_0 + b_1 x_1 + b_2 x_2 + \dots + b_n x_n$$


Polynomial
Linear
Regression

$$y = b_0 + b_1 x_1 + b_2 x_1^2 + \dots + b_n x_1^n$$


Simple Linear Regression




Simple Linear Regression



Polynomial Regression



Polynomial Regression



Polynomial Regression

Polynomial
Linear
Regression

$$y = b_0 + b_1x_1 + b_2x_1^2 + \dots + b_nx_1^n$$

```
data_preprocessing_template.py polynomial_regression.py*
16
17# Feature Scaling
18"""from sklearn.preprocessing import StandardScaler
19sc_X = StandardScaler()
20X_train = sc_X.fit_transform(X_train)
21X_test = sc_X.transform(X_test)"""
22
23# Fitting Linear Regression to the dataset
24from sklearn.linear_model import LinearRegression
25lin_reg = LinearRegression()
26lin_reg.fit(X, y)
27
28# Fitting Polynomial Regression to the dataset
29from sklearn.preprocessing import PolynomialFeatures
30poly_reg = PolynomialFeatures(degree = 4)
31X_poly = poly_reg.fit_transform(X)
32poly_reg.fit(X_poly, y)
33lin_reg_2 = LinearRegression()
34lin_reg_2.fit(X_poly, y)
35
36# Visualising the Linear Regression results
37plt.scatter(X, y, color = 'red')
38plt.plot(X, lin_reg.predict(X), color = 'blue')
39plt.title('Truth or Bluff (Linear Regression)')
40plt.xlabel('Position level')
41plt.ylabel('Salary')
42plt.show()
43
44# Visualising the Polynomial Regression results
45plt.scatter(X, y, color = 'red')
46plt.plot(X, lin_reg_2.predict(poly_reg.fit_transform(X)), color = 'blue')
47plt.title('Truth or Bluff (Polynomial Regression)')
48plt.xlabel('Position level')
49plt.ylabel('Salary')
50plt.show()
51
52# Predicting a new result with Linear Regression
53lin_reg.predict(6.5)
54
55# Predicting a new result with Polynomial Regression
56lin_reg_2.predict(poly_reg.fit_transform(6.5))
```

Simple Linear Regression Model

Polynomial Regression Model

Simple Linear Regression Model Prediction

Polynomial Regression Model Prediction

Name	Type	Size	Value
X	int64	(10, 1)	array([[1], [2], [3], [4], [5], [6], [7], [8], [9], [10]])
X_poly	float64	(10, 5)	array([[1.00000000e+00, 1.00000000e+00, 1.00000000e+00, 1.00000000e+00, 1.00000000e+00], [1.62890563e+00, 1.62890563e+00, 1.62890563e+00, 1.62890563e+00, 1.62890563e+00], [2.77555736e+00, 2.77555736e+00, 2.77555736e+00, 2.77555736e+00, 2.77555736e+00], [4.75246284e+00, 4.75246284e+00, 4.75246284e+00, 4.75246284e+00, 4.75246284e+00], [7.58142943e+00, 7.58142943e+00, 7.58142943e+00, 7.58142943e+00, 7.58142943e+00], [11.22472109e+00, 11.22472109e+00, 11.22472109e+00, 11.22472109e+00, 11.22472109e+00], [15.74132969e+00, 15.74132969e+00, 15.74132969e+00, 15.74132969e+00, 15.74132969e+00], [21.09026521e+00, 21.09026521e+00, 21.09026521e+00, 21.09026521e+00, 21.09026521e+00]])
dataset	DataFrame	(10, 3)	Column names: Position, Level, Salary
y	int64	(10,)	array([45000, 50000, 60000, 80000, 110000, 150000, 200000, 300000, 500000, 1000000])

Object inspector Variable explorer File explorer

IPython console

Console 1/A

```
...: plt.ylabel('Salary')
...: plt.show()

In [14]: lin_reg.predict(6.5)
Out[14]: array([ 330378.78787879])


In [15]: plt.scatter(X, y, color = 'red')
...: plt.plot(X, lin_reg.predict(X), color = 'blue')
...: plt.title('Truth or Bluff (Linear Regression)')
...: plt.xlabel('Position level')
...: plt.ylabel('Salary')
...: plt.show()

In [16]: lin_reg_2.predict(poly_reg.fit_transform(6.5))
Out[16]: array([ 1158862.45265153])

In [17]:
```

Permissions: RW End-of-lines: LF Encoding: UTF-8-GUESSED Line: 56 Column: 47 Memory: 38 %

Simple Linear Regression



Polynomial Regression

